

Fast Error-Bounded HPC Data Compressor (sz-1.1)

User Guide (Version 1.1)

Mathematics and Computer Science (MCS)

Argonne National Laboratory

Contact: Sheng Di (sdi1@anl.gov)

June, 2016

Table of Contents

Table of Contents	1
1. Brief description.....	1
2. How to install SZ	2
3. Quick Start	3
3.1 Compression	3
3.2 Decompression	4
4. Application Programming Interface (API)	4
4.1 Compression/Decompression by C Interfaces.....	5
4.2 Compression/Decompression by Fortran Interfaces.....	7
5 Test cases.....	10
6. Version history	11
7. Q&A and Trouble shooting.....	12

1. Brief description

- **SZ** (short for **Squeeze**) is an **error-bounded** HPC in-situ data compressor for significantly reducing the data sizes, which can be leveraged to improve the checkpoint/restart performance and post-processing efficiency for HPC executions.
- SZ can be used to compress different types of data (single-precision and double-precision) and any shapes of the array. Current version supports up to five dimensions. Higher dimensions can also be extended easily.
- SZ is very easy to use. It supports three programming languages: Fortran, C and Java.
- SZ supports many different architectures, including x86_32bits (denoted by linux_x86 in the Makefile), x86_64bits (denoted by linux_x64 in the Makefile), ARM (denoted by linux_arm), SOLARIS (denoted by solaris), IBM BlueGene series (denoted by pps).

- SZ allows setting the compression error bound based on *absolute error bound* and/or *relative error bound*, by using `sz.config` (which can be found in the directory *example*) or by passing arguments through programming interfaces.
 - **Absolute error bound** (namely *absErrBound* in the configuration file `sz.config`): It is to limit the (de)compression errors to be within an absolute error. For example, `absErrBound=0.0001` means the decompressed value must be in $[V-0.0001, V+0.0001]$, where V is the original true value.
 - **Relative error bound** (called *relBoundRatio* in the configuration file `sz.config`): It is to limit the (de)compression errors by considering the global data value range size (i.e., taking into account the range size (`max_value - min_value`)). For example, suppose `relBoundRatio` is set to 0.01, and the data set is $\{100, 101, 102, 103, 104, \dots, 110\}$. That is, the maximum value is 110 and minimum value is 100. So, the global value range size is $110-100=10$, and the error bound will actually be $10 \times 0.01 = 0.1$, from the perspective of "relBoundRatio".
- Users can set the real compression error bound based on only `absErrorBound`, `relBoundRatio`, or a kind of combination of them. Two types of combinations are provided: **AND**, **OR**. **ABS_AND_REL** means that both of the two bounds (`absErrorBound` and `relBoundRatio`) will be considered in the compression. **ABS_OR_REL** means that the compression error is satisfied as long as one type of bound is met.
- If there are many variables to be compressed, we recommend to compress them using batch-compression way. Specifically, there are two steps in the batch-compression: (1) register/add variables, and (2) perform the compression. Please reference the description of `SZ_batchAddVar()` and `SZ_batch_compress()`. An example code (`testfloat_batch_compress.c`) can also be found in `example/` directory.
- Users are allowed to set the endian type of the data in the `sz.config`. Please check the comments of this file in the `example/` directory.

2. How to install SZ

The SZ software can be downloaded from <http://collab.mcs.anl.gov/display/ESR/SZ>

Perform the following three simple steps to finish the installation:

```
configure --prefix=[INSTALL_DIR]
make
make install
```

Then, you'll find all the executables in `[INSTALL_DIR]/bin` and `.a` and `.so` libraries in `[INSTALL_DIR]/lib`

3. Quick Start

The testing cases can be found in **[SZ_Package]/example**

You can use "make clean;make" to recompile all the example codes, or compile them by the customized Makefile.bk as follows:

```
make -f Makefile.bk
```

(Makefile.bk allows you to compile your customized source codes.)

3.1 Compression

Testing commands:

Run **`./testdouble_compress sz.config testdouble_8_8_128.dat 8 8 128`** to compress the data `testdouble_8_8_128.dat`.

Run **`./testdouble_compress sz.config testdouble_8_8_8_128.dat 8 8 8 128`** to compress the data `testdouble_8_8_8_128.dat`.

Run **`./testfloat_compress sz.config testfloat_8_8_128.dat 8 8 128`** to compress the data `testfloat_8_8_128.dat`

Remark:

`testdouble_8_8_128.dat` and `testdouble_8_8_8_128.dat` are two binary testing files, which contain a 3d array (128X8X8) and a 4d array (128X8X8X8) respectively. Their data values are shown in the two plain text files, `testdouble_8_8_128.txt` and `testdouble_8_8_8_128.txt`. These two data files are from FLASH_Blast2 and FLASH_MacLaurin respectively (the two test data are both extracted at time step 100). The compressed data files to be generated are named `testdouble_8_8_128.dat.sz` and `testdouble_8_8_8_128.dat.sz` respectively.

`./testfloat_compress.c` is an example to show how to compress single-precision data. Use `testfloat_8_8_128.dat` as the input when testing the compression of single-precision data.

`sz.config` is the configuration file. The key settings are *errorBoundMode*, *absErrBound*, and *relBoundRatio*, which are described below.

- ***absErrBound*** refers to the absolute error bound, which is to limit the (de)compression errors to be within an absolute error. For example, `absErrBound=0.0001` means the decompressed value must be in $[V-0.0001, V+0.0001]$, where V is the original true value.
- ***relBoundRatio*** refers to relative bound ratio, which is to limit the (de)compression errors by considering the global data value range size (i.e., taking into account the range size (`max_value - min_value`)). For example, suppose `relBoundRatio` is set to 0.01, and the data set is {100,101,102,103,104,...,110}. In this case, the maximum value is 110 and the minimum is 100. So, the global value range size is $110-100=10$, and the error bound will be $10*0.01=0.1$, from the perspective of "relBoundRatio".

- **errorBoundMode** is to define a combination of the above two types of error bounds. There are four types of values: **ABS**, **REL**, **ABS_AND_REL**, **ABS_OR_REL**.
 - **ABS** takes only "absolute error bound" into account. That is, relative bound ratio will be ignored.
 - **REL** takes only "relative bound ratio" into account. That is, absolute error bound will be ignored.
 - **ABS_AND_REL** takes both of the two bounds into account. The compression errors will be limited using both `absErrBound` and `relBoundRatio*rangeSize`. That is, the two bounds must be both met.
 - **ABS_OR_REL** takes both of the two bounds into account. The compression errors will be limited using either `absErrBound` or `relBoundRatio*rangeSize`. That is, only one bound is required to be met.

sz.config is the configuration file used to set the compression environment. Please read the comment in the file to understand the parameters.

3.2 Decompression

Testing commands:

```
./testdouble_decompress sz.config testdouble_8_8_128.dat.sz
./testdouble_decompress sz.config testdouble_8_8_8_128.dat.sz
./testfloat_decompress sz.config testfloat_8_8_128.dat.sz
```

Remark:

- Unlike compression, you don't have to provide the error bound information (such as `errBoundMode`, `absErrBound`, and `relBoundRatio`), when performing the data decompression, because such information is stored in the compressed data stream.
- The output files of the `test_decompress.c` are `.out` files, i.e., `testdouble_8_8_128.dat.sz.out` and `testdouble_8_8_8_128.dat.sz.out` respectively. You can compare `.txt` file and `.out` file for checking the compression errors for each data point. For instance, compare `testdouble_8_8_8_128.txt` and `testdouble_8_8_8_128.dat.sz.out`.

4. Application Programming Interface (API)

Programming interfaces are provided in two programming languages – C and Fortran (SZ-0.x versions also provided Java interfaces). The usage methods of the interfaces are quite similar across different programming languages, with only a few differences. For example, In C interface, a *dataType* (either `SZ_FLOAT` or `SZ_DOUBLE`) is required, while Fortran interface does not require this argument because of the function overloading feature.

4.1 Compression/Decompression by C Interfaces

There are three key interfaces for compression/decompression in C.

- (1) Initialize the compressor by calling `SZ_Init()`;
- (2) Compress the data (a floating-point array) by `SZ_compress()`, or decompress the data by `SZ_decompress()`;
- (3) Finalize the compressor by `SZ_Finalize()` if the compressor won't be used any more.

Interfaces:

(a) **SZ_Init**

Initialize the SZ compressor. `SZ_Init()` just needs to be called only **once** before performing multiple compressions for different variables (data arrays).

Synopsis: **void** `SZ_Init`(**char** *`configFilePath`);

Input:

configFilePath the configuration file path (such as `example/sz.config`)

Return: 1 (failure) or 0 (success)

(b) **SZ_compress**

Compress the floating-point data array. Two types of interfaces are provided, as shown below. For the first one, the three important control parameters (`errBoundMode`, `absErrBound`, and `relBoundRatio`) will be given by the configuration file `sz.config`. For the second one, the three control parameters will be passed using arguments, so in this case, the parameter settings in the `sz.config` will be ignored.

There are three compression interfaces with different arguments, as listed below. The user just needs to choose one of them in compressing data.

Synopsis:

char *`SZ_compress`(**int** `dataType`, **void** *`data`, **int** *`outSize`, **int** `r5`, **int** `r4`, **int** `r3`, **int** `r2`, **int** `r1`);

char *`SZ_compress_args`(**int** `dataType`, **void** *`data`, **int** *`outSize`, **int** `errBoundMode`, **double** `absErrBound`, **double** `relBoundRatio`, **int** `r5`, **int** `r4`, **int** `r3`, **int** `r2`, **int** `r1`);

int `SZ_compress_args2`(**int** `dataType`, **void** *`data`, **char*** `compressed_bytes`, **int** *`outSize`, **int** `errBoundMode`, **double** `absErrBound`, **double** `relBoundRatio`, **int** `r5`, **int** `r4`, **int** `r3`, **int** `r2`, **int** `r1`);

Input:

dataType the indicator that indicates the data type
(two options: either `SZ_FLOAT` or `SZ_DOUBLE`)

data the variable that contains the data to be compressed.
(Current version only supports “double precision” data)

compressed_bytes the address that contains the compressed bytes

outSize	the data stream size (in bytes) after compression.
r5	size of dimension 5
r4	size of dimension 4
r3	size of dimension 3
r2	size of dimension 2
r1	size of dimension 1

Return: Compressed data stream (in the form of bytes)

Usage tips: The dimension of the variable is determined based on the five dimension parameters (r5, r4, r3, r2, and r1). For instance, if the variable is a 2D array (M X N), then r5=0, r4=0, r3=0, r2=M, and r1=N. If the variable to protect is a 4D array, then only r5 is set to 0. (See test_compress.c for details).

(c) SZ_decompress

Decompress/recover the data. Two options, as listed below.

Synopsis:

```
void *SZ_decompress(int dataType, char *bytes, int byteLength,
                    int r5, int r4, int r3, int r2, int r1);
int SZ_decompress_args(int dataType, char *bytes, int byteLength,
                       void* decompressed_array,
                       int r5, int r4, int r3, int r2, int r1);
```

Input:

dataType	the indicator to indicate the data type (either <i>SZ_FLOAT</i> or <i>SZ_DOUBLE</i>)
bytes	the compressed data stream to be decompressed
byteLength	length of the compressed data stream
decompressed_array	the address to store decompressed data
r5	size of dimension 5
r4	size of dimension 4
r3	size of dimension 3
r2	size of dimension 2
r1	size of dimension 1

Return: the recovered data array decompressed from the compressed bytes.

(d) SZ_batchAddVar

Register/add a variable (denoted by *var*) to be compressed with other variables together in a batch way.

Synopsis:

```
void SZ_batchAddVar(char* varName, int dataType, void* var,
                    int r5, int r4, int r3, int r2, int r1,
                    int errBoundMode, double absErrBound, double relBoundRatio);
```

(e) SZ_batchDelVar

Deregister/delete a variable (denoted by *var*) from the list of registered variables, that are to be compressed with other variables together in a batch way.

Synopsis:

```
int SZ_batchDelVar(char* varName);
```

Input:

varName the name of variable used in the registration.

Return: 0: success or 1: no corresponding variable is found based on varName.

(f) SZ_batch_compress

Compress the data in a batch way: all of the registered variable data will be compressed together (The benefit is improvement of compression factor).

Synopsis:

char* SZ_batch_compress(**int** *outSize);

Input:

outSize the data stream size (in bytes) after compression.

Return: the compressed stream.

(g) SZ_batch_decompress

Decompress the batch-compressed stream.

Synopsis:

SZ_VarSet* SZ_batch_decompress (**char*** compressedStream,
int compressedLength);

Input:

compressedStream the compressed stream

compressedLength the length of the compressed stream (in byte)

Return: The data structure containing the decompressed data with multiple variables.

See VarSet.h for more details. The global SZ_VarSet is defined in sz.h: SZ_VarSet* sz_varset.

~~(h) SZ_Finalize~~

Release the memory and compression environment.

This function is **deprecated** in sz 1.0 or later versions.

Synopsis: **int** SZ_Finalize();

Input: none.

Return: none.

4.2 Compression/Decompression by Fortran Interfaces

Interfaces:

(a) SZ_Init

Initialize the SZ compressor. SZ_Init() just needs to be called only **once** before performing multiple compressions for different variables (data arrays).

Synopsis: **SZ_Init**(configFilePath, ierr);

Input:

configFilePath configuration file path (e.g., sz.config)

CHARACTER(len=32) :: configFilePath

Output:

ierr successful (0) or failed (1)

INTEGER(Kind=4) :: ierr

(b) SZ_Compress

Compress the floating-point data array. Two types of interfaces are provided, as shown below. For the first one, the three important control parameters (`errBoundMode`, `absErrBound`, and `relBoundRatio`) will be given by the configuration file `sz.config`. For the second one, the three control parameters will be passed using arguments, so in this case, the parameter settings in the `sz.config` will be ignored.

Synopsis A:

SZ_compress(`data`, `bytes`, `outSize`);

Input:

data the data array to be compressed
(the data here is a floating-point data array with up to 5 dimensions. For example, "REAL(KIND=8), DIMENSION(:, :, :) :: *data*" indicates a 3D double-precision array, where *data* refers to the array variable.)

Output:

bytes the byte stream generated after the compression
INTEGER(kind=1), DIMENSION(:), allocatable :: bytes

outsize the size (in bytes) of the byte stream
INTEGER(kind=4) :: OutSize

Synopsis B:

SZ_Compress (`data`, `bytes`, `outSize`,
`errBoundMode`, `absErrBound`, `relBoundRatio`);

Input:

data the data array to be compressed
(the data here is a floating-point data array with up to 5 dimensions. For example, "REAL(KIND=8), DIMENSION(:, :, :) :: *data*" indicates a 3D double-precision array, where *data* refers to the array variable.)

errBoundMode the error bound mode.
Four options: ABS, REL, ABS_AND_REL, ABS_OR_REL
INTEGER(kind=4) :: ErrBoundMode

absErrBound absolute error bound
REAL(kind=4 or 8) :: absErrBound

relBoundRatio relative bound ratio
REAL(kind=4 or 8) :: relBoundRatio
(Details about error bound mode, absolute error bound, and relative bound ratio can be found in Section 3.1)

Output:

bytes the byte stream generated after the compression
INTEGER(kind=1), DIMENSION(:), allocatable :: bytes

outsize the size (in bytes) of the byte stream
INTEGER(kind=4) :: OutSize

(c) **SZ_Decompress**

Decompress/recover the data

Synopsis:

SZ_Decompress(bytes, data, [r1,r2,...])

Input:

bytes	the compressed data stream to be decompressed INTEGER(kind=1), DIMENSION(:) :: Bytes
data	length of the compressed data stream REAL(KIND=4 or 8), DIMENSION(:, :, ..., :), allocatable :: data
r1	size of dimension 1
r2	size of dimension 2
r3	size of dimension 3
r4	size of dimension 4
r5	size of dimension 5 INTEGER(kind=4) :: r1[, r2, r3, r4, r5]

Usage tips: SZ_Decompress supports the decompression of the array with at most 5 dimensions. The dimension sizes (such as r1, r2,) are supposed to be provided. For example, in order to decompress a binary stream whose original data is a 3D array (r3=10,r2=8,r1=8), the function is like "SZ_Decompress(bytes, data, 8, 8, 10).

(d) **SZ_BatchAddVar**

Register/add a data variable (denoted by *var*) to be compressed with other variables together in a batch way.

Synopsis:

void SZ_batchAddVar(varName, var,
errBoundMode, absErrBound, relBoundRatio);

varName	the name of the variable to be registered/added CHARACTER(len=128) :: varName
var	the variable/data to be registered/added
errBoundMode	the error bound mode. Four options: ABS, REL, ABS_AND_REL, ABS_OR_REL INTEGER(kind=4) :: ErrBoundMode
absErrBound	absolute error bound REAL(kind=4 or 8) :: absErrBound
relBoundRatio	relative bound ratio REAL(kind=4 or 8) :: relBoundRatio (Details about error bound mode, absolute error bound, and relative bound ratio can be found in Section 3.1.

(e) **SZ_BatchDelVar**

Deregister/delete a variable (denoted by *var*) from the list of registered variables, that are to be compressed with other variables together in a batch way.

Synopsis:

void SZ_batchDelVar(varName, ierr);

Input:

varName the name of variable used in the registration.

CHARACTER(len=128) :: varName

Output:

ierr the output status (0: success or 1: no variable found)

INTEGER(kind=4) :: ErrBoundMode

Return: 0: success or 1: no corresponding variable is found based on varName.

(f) SZ_Batch_Compress

Compress the data in a batch way: all of the registered variable data will be compressed together (The benefit is improvement of compression factor).

Synopsis:

void SZ_Batch_Compress(bytes, outSize)

Output:

bytes the byte stream generated after the compression

INTEGER(kind=1), DIMENSION(:), allocatable :: bytes

outsize the size (in bytes) of the byte stream

(g) SZ_Batch-Decompress

Decompress the batch-compressed stream.

Synopsis:

void SZ_Batch-Decompress(bytes, outSize)

Output:

bytes the compressed data stream to be decompressed

INTEGER(kind=1), DIMENSION(:) :: Bytes

outsize the size of the decompressed data stream

INTEGER(kind=4) :: OutSize

(h) SZ_Finalize

Release the memory and compression environment

Synopsis: **SZ_Finalize();**

Input: none.

Return: none.

5 Test cases

example/testdouble_compress.c

example/testdouble_decompress.c

example/testfloat_compress.c

example/testfloat_decompress.c

example/testfloat_batch_compress.c

example/testdouble_batch_compress.c

example/testdouble_compress.f90

example/testdouble/decompress.f90

6. Version history

The latest version (**version 1.1**) is the recommended one.

version New features

SZ 0.2-0.4 Compression ratio is the same as SZ 0.5. The key difference is different implementation ways, such that SZ 0.5 is much faster than SZ 0.2-0.4.

SZ 0.5.1 Support version checking

SZ 0.5.2 finer compression granularity for unpredictable data, and also remove redundant Java storage bytes

SZ 0.5.3 Integrate with the dynamic segmentation support

SZ 0.5.4 Gzip_mode: default --> fast_mode ; Support reserved value

SZ 0.5.5 runtime memory is shrinked (by changing int xxx to byte xxx in the codes)

The bug that writing decompressed data may encounter exceptions is fixed.

Memory leaking bug for ppc architecture is fixed.

SZ 0.5.6 improve compression ratio for some cases (when the values in some segmentation are always the same, this segment will be merged forward)

SZ 0.5.7 improve the decompression speed for some cases

SZ 0.5.8 Refine the leading-zero granularity (change it from byte to bits based on the distribution). For example, in SZ0.5.7, the leading-zero is always in bytes, 0, 1, 2, or 3. In

SZ0.5.8 The leading-zero part could be xxxx xxxx xx xx xx xxxx xxxx (where each x means a bit in the leading-zero part)

SZ 0.5.9 optimize the offset by using simple right-shifting method. Experiments show that this cannot improve compression ratio actually, because simple right-shifting actually make each data be multiplied by $2^{\{-k\}}$, where k is # right-shifting bits. The pros is to save bits because of more leading-zero bytes, but the cons is much more required bits to save. A good solution is SZ 0.5.10!

SZ 0.5.10 optimize the offset by using the optimized formula of computing the median_value based on optimized right-shifting method. Anyway, SZ0.5.10 improves compression ratio a lot for hard-to-compress datasets. (Hard-to-compress datasets refer to the cases whose compression ratios are usually very limited)

SZ 0.5.11 In a very few cases, SZ 0.5.10 cannot guarantee the error-bounds to a certain user-specified level. For example, when absolute error bound = $1E-6$, the maximum decompression error may be 0.01(>> $1E-6$) because of the huge value range even in the optimized segments such that the normalized data cannot reach the required precision even soaring all of the 64 or 32 mantissa bits. SZ 0.5.11 fixed the problem well, with degraded compression ratio less than 1%.

SZ 0.5.12 A parameter setting called "offset" is added to the configuration file sz.config. The value of offset is an integer in [1,7]. Generally, we recommend offset=2 or 3, while we also find that some other settings (such as offset=7) may lead to better compression ratios in some cases. How to automatize/optimize the selection of offset value would be the future work. In addition, the compression speed is improved, by replacing java List by array

implementation in the code.

SZ 0.5.13 Compression performance is improved, by replacing some class instances in the source code by primitive data type implementation.

SZ 0.5.14 fixed a design bug, which improves the compression ratio further.

SZ 0.5.15 improved the compression ratio for single-precision data compression, by tuning the offset.

The version 0.x were all coded in Java, and C/Fortran interfaces were provided by using JNI and C/Fortran wrapper. SZ 1.0 is coded in C purely.

SZ 1.0 Pure C version. In this version, the users don't need to install JDK and make the relative configurations any more. It provides dataEndianType in the sz.config file, so it can be used to compress the data file which was generated on different endian-type systems.

SZ 1.1 batch_compression function is added to this version. Compression performance is improved slightly due to for(;;) being replaced by memcpy() somewhere.

7. Q&A and Trouble shooting

1. Do I need to call SZ_init() every time I compress a variable in the program?

Answer: No. In the progress, SZ_init() just needs to be called once at the beginning, and thereafter you can always compress different variables using the compression/decompression functions on demand, until SZ_finalize() is called.

2. If I want to use SZ_compress_args() function and specify the errorBoundMode and bounds at run time instead of using the sz.config, do I need to call SZ_init()?

Answer: It depends. In fact, sz.config has some important parameter settings, e.g., data_endian_type (little or big). You can also set these parameters manually in your code or use the default setting. Please check sz.h and conf.c for details. We highly recommend to use sz.config to initialize the compression environment, because some critical parameters such as dataEndianType may be random numbers without such an initialization.

3. How to deal with “Error: The input file or data stream is not in SZ format!”?

Answer: This error is because the input file or data stream used to be decompressed is probably not the byte stream compressed/generated by the SZ. Please use the compressed file (such as data.sz) in the decompression.

<END>